

Methods 3: Multilevel Statistical Modeling and Machine Learning

Week 7: *Linear regression revisited (machine learning)*

November 9, 2021

by: Lau Møller Andersen

These slides are distributed according to the CC
BY 4.0 licence:

<https://creativecommons.org/licenses/by/4.0/>



Attribution 4.0 International (CC BY 4.0)

Last time – Mid-way evaluation

- 1) Write something you liked about the course so far
- 2) Write something you did not like about the course so far
- 3) What would you change?

Summary

summary of student feedback

LIKE

- Good amount of practical exercises (n=3)
- Hands-on coding (n=8)
- GitHub introduction (n=3)
- Lau is answering questions (n=9)
- That we go into depth (n=1)
- Peer-review (n=1)
- Good readings (n=5)
- Involvement of students (n=2)
- Lectures are interesting (n=2)
- Good structure of classes (n=4)
- Lau is engaged (n=4)
- Exercises tied with papers (n=1)
- Allow more time for clarification (n=1)
- Easy to get help (n=3)
- Pace too slow (n=1)
- Explaining lecture goals (n=1)
- Demystifying math (n=2)

DID NOT LIKE

- Portfolio too challenging or too many questions (n=8)
- GitHub is not worth it (n=1)
- Math is hard - Explain math more (n=5)
- Coding is hard (n=2)
- Too many slides, not enough time (n=4)
- Questions not clear enough in portfolio (n=9)
- Connection between lecture and readings (n=5)
- Bad structure of classes (n=1)
- Exercises that have not been explicitly covered in class (n=5)
- Not clear why we are doing the assignments (n=1)
- That Lau is not in all classes (n=2)
- Methods 1, 2 and 3 are not well connected (n=2)
- Lau talking too fast and not loud enough (n=1)

CHANGE

- Closer match between lectures and classes (n=2)
- List main concepts in syllabus (n=1)
- Better deadlines needed (n=5)
- More explanation on portfolios (n=1)
- More feedback on exercises (n=3)
- Programming together (n=1)
- Go through papers that exercises are based on (n=1)

Total n = 29

https://github.com/ualsbombe/github_methods_3/blob/main/week_07/mid-way_evaluation.pdf

Summary – what you liked

- Top 3
 - 1) Lau is answering questions (n=9)
 - 2) Good readings (n=5)
 - 3) Lau is engaged (n=4) Good structure of classes (n=4)

Total n = 29

Summary – what you did not like

- Top 3
 - 1) Questions not clear enough in portfolio (n=9)
 - 2) Portfolio too challenging or too many questions (n=8)
 - 3) Math is hard – explain math more (n=5)
Connection between lecture and readings (n=5)
Exercises not explicitly covered in class (n=5)

Total n = 29

Summary – to change

CHANGE

Closer match between lectures and classes (n=2)

List main concepts in syllabus (n=1)

Better deadlines needed (n=5)

More explanation on portfolios (n=1)

More feedback on exercises (n=3)

Programming together (n=1)

Go through papers that exercises are based on (n=1)

Summary – what I promise to change

I will be very careful and aim at writing questions that are easy to understand

Exercise for you: In tomorrow's exercise: For each question indicate whether you understood what was required of you.

Summary – what I promise to change

Connection between lectures and class will be closer – we will be following the textbook more closely

Summary – a discussion

Lau is answering questions (n=9)

Too many slides, not enough time (n=4)

This is a classic conundrum (which we should still try to solve)

Summary – a discussion

What is the optimal balance between me going through the slides and me answering questions?

Discuss for 3-5 minutes

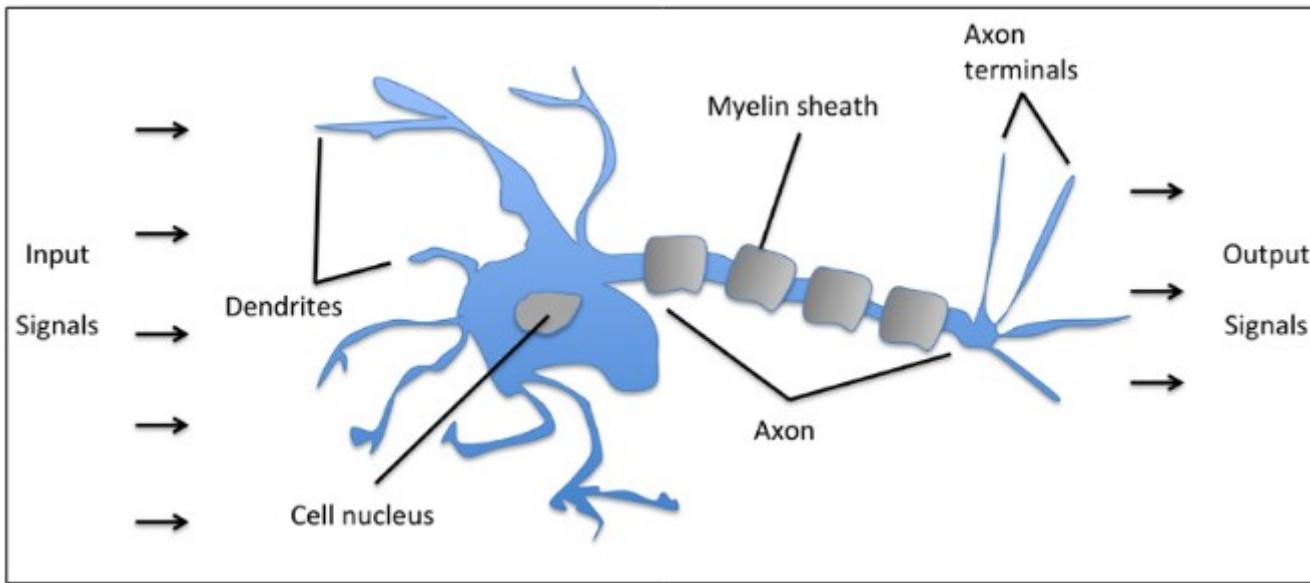
What should our balance be?

Learning goals

Linear regression revisited (machine learning)

- 1) Learning some early classification methods
perceptron and adaline
- 2) Learning how linear regression (with biasing penalties) can be constructed and cross-validated
- 3) Understanding that biasing in-sample solutions can improve out-of-sample predictions
biasing (tradeoff between variance and bias) conflicts with OLS (minimizing RSS)

Black box idea



$$\mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$$

$$z = w_1 x_1 + \dots + w_m x_m$$

Question: what do x , w and z correspond to in the above picture of the *Perceptron*?

(p. 18: Raschka, 2015)

Prediction/classification rule

like a neuron, either it fires or then it does not

det her symbol: er either or

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq \theta^{\text{theta}} \\ -1 & \text{otherwise} \end{cases}$$

Perceptron fires

Perceptron doesn't fire

θ is a pre-specified threshold

(Raschka, 2015)

Prediction/classification rule

intercept som minus beta

$$w_0 = -\theta$$

$$z = w_0 x_0 + w_1 x_1 + \dots + w_m x_m = \mathbf{w}^T \mathbf{x}$$

$$x_0 = 1$$

$$z = -\theta + w_1 x_1 + \dots + w_m x_m = \mathbf{w}^T \mathbf{x}$$

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

Perceptron fires
Perceptron doesn't fire

(Raschka, 2015)

Perceptron classification

Bonus info

fra portfolio

$$f(x) = a + \frac{b-a}{1+e^{\frac{c-x}{d}}}$$

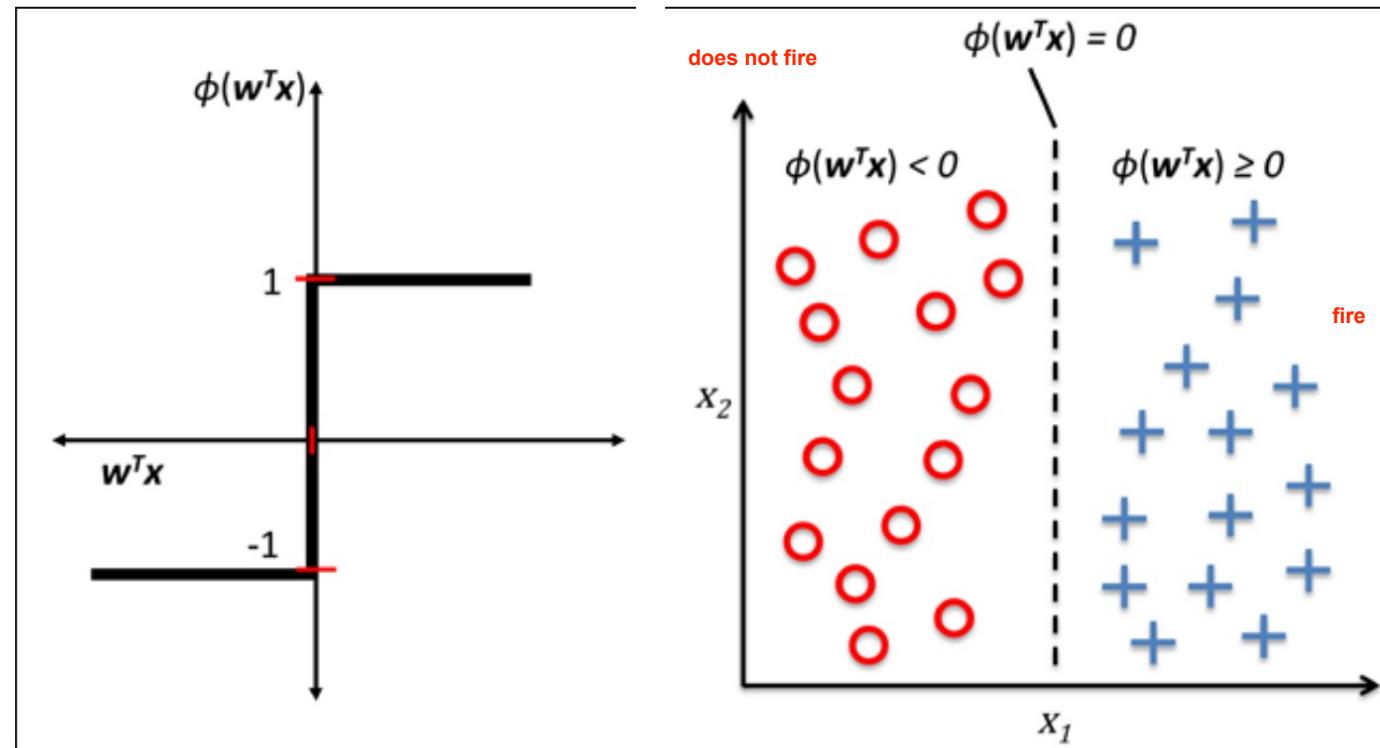
Practical
exercise 5...

$$a = -1$$

$$b = 1$$

$$c = 0$$

d going to 0



(p. 21: Raschka, 2015)

We want to find $w^T x$ that achieves this separation

In Python

```
class Perceptron(object):
    """ Perceptron classifier

    Parameters
    -----
    eta : float
        Learning rate (between 0.0 and 1.0)
    n_iter : int
        Passes over the training dataset.

    Attributes
    -----
    w_ : 1d-array
        Weights after fitting.
    errors_ : list
        Number of misclassifications in every epoch.

    """
    alt med under_betyder at noget er blevet fittet
```

Special definition that indicates what the object (*Perceptron*) can be initialised with

```
dobbelts__ def __init__(self, eta=0.01, n_iter=10):
    self.eta = eta
    self.n_iter = n_iter
```

```
ppn = Perceptron(eta=0.1, n_iter=10)
```

Specifying methods of Perceptron

this perceptron fit function will find that line (the step function line)

1. Initialize the weights to 0 or small random numbers.
2. For each training sample $x^{(i)}$ perform the following steps:
 1. Compute the output value \hat{y} .
 2. Update the weights.

```
self refers to the variable within the class
def fit(self, X, y):
    """ Fit training data.

    Parameters
    -----
    X : {array-like}, shape = [n_samples, n_features]
        Traing vectors, where n_samples
        is the number of samples and
        n_features is the number of features.
    y : array-like, shape = [n_samples]
        Target values.

    Returns
    -----
    self : object
    """
    self.w_ = np.zeros(1 + X.shape[1])
    self.errors_ = []

    for _ in range(self.n_iter):
        errors = 0
        for xi, target in zip(X, y):
            update = self.eta * (target - self.predict(xi))
            self.w_[1:] += update * xi
            self.w_[0] += update
            errors += int(update != 0.0)
        self.errors_.append(errors)
    return self
```

Compute the output value \hat{y}

```
def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_[1:]) + self.w_[0]
                                         X*BETA
                                         constanten
def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.net_input(X) >= 0.0, 1, -1)
                                         if else her :))
```

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ -1 & \text{otherwise} \end{cases}$$

When we are right

$$\Delta w_j = \eta \left(\begin{matrix} \text{real label} \\ -1 \end{matrix} - \begin{matrix} \text{predicted label} \\ -1 \end{matrix} \right) x_j^{(i)} = 0$$

0=weight does not change (we were right)

$$\Delta w_j = \eta \left(\begin{matrix} \text{real label} \\ 1 \end{matrix} - \begin{matrix} \text{predicted label} \\ 1 \end{matrix} \right) x_j^{(i)} = 0$$

0=weight does not change (we were right)

Δw_j : change in weight

η : learning rate

When we are wrong

when we are wrong : change the learning rate

real label

$$\Delta w_j = \eta (1 - 1) x_j^{(i)} = \eta (2) \underline{x_j^{(i)}}$$

predicted label

here we have to adjust the weight

real label

$$\Delta w_j = \eta (-1 - 1) x_j^{(i)} = \eta (-2) \underline{x_j^{(i)}}$$

predicted label

here we have to adjust the weight

Δw_j : change in weight

η : learning rate

the learning rate is something we set ourselves:

- how do we define it?
- is there a convention?

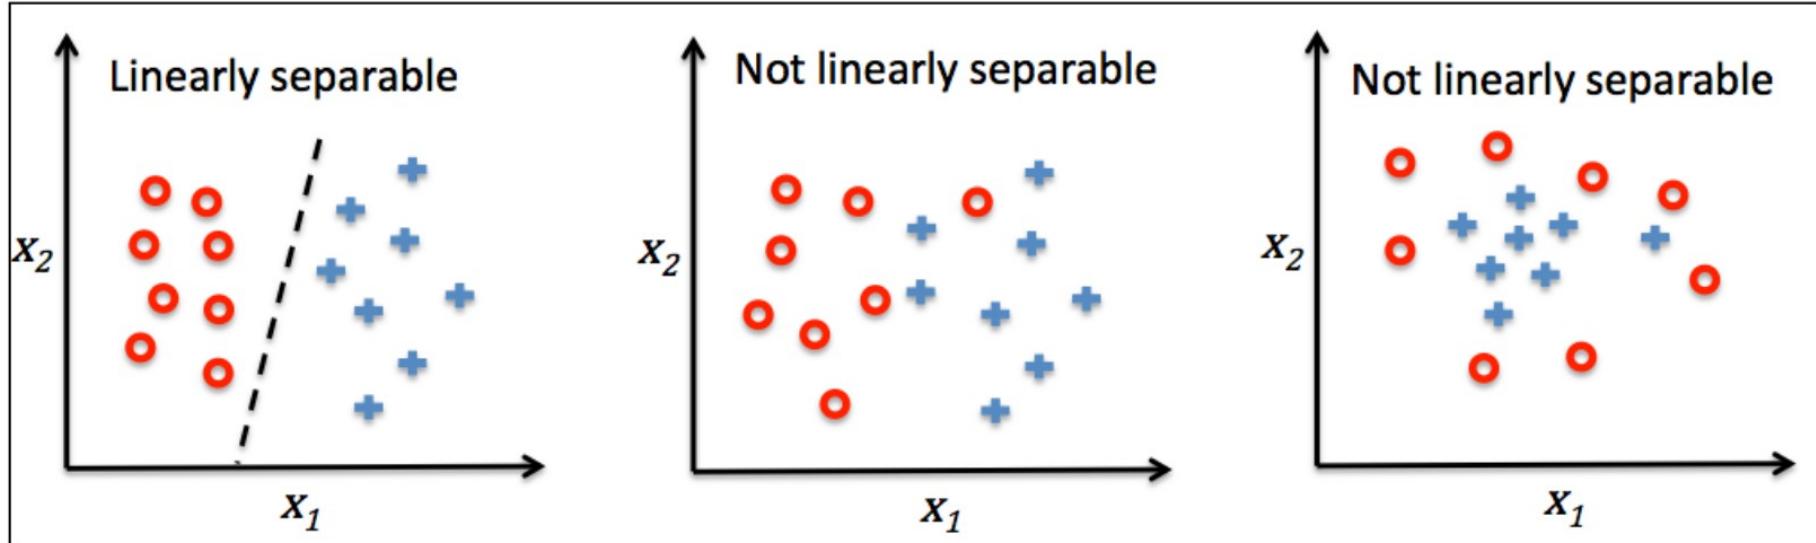
depends on the range of your x, in laus example it is set to 0.1..

Men ærligt forstår jeg ikke hvordan den skal defineres.

For learning rate må ikke være for høj eller for lav..

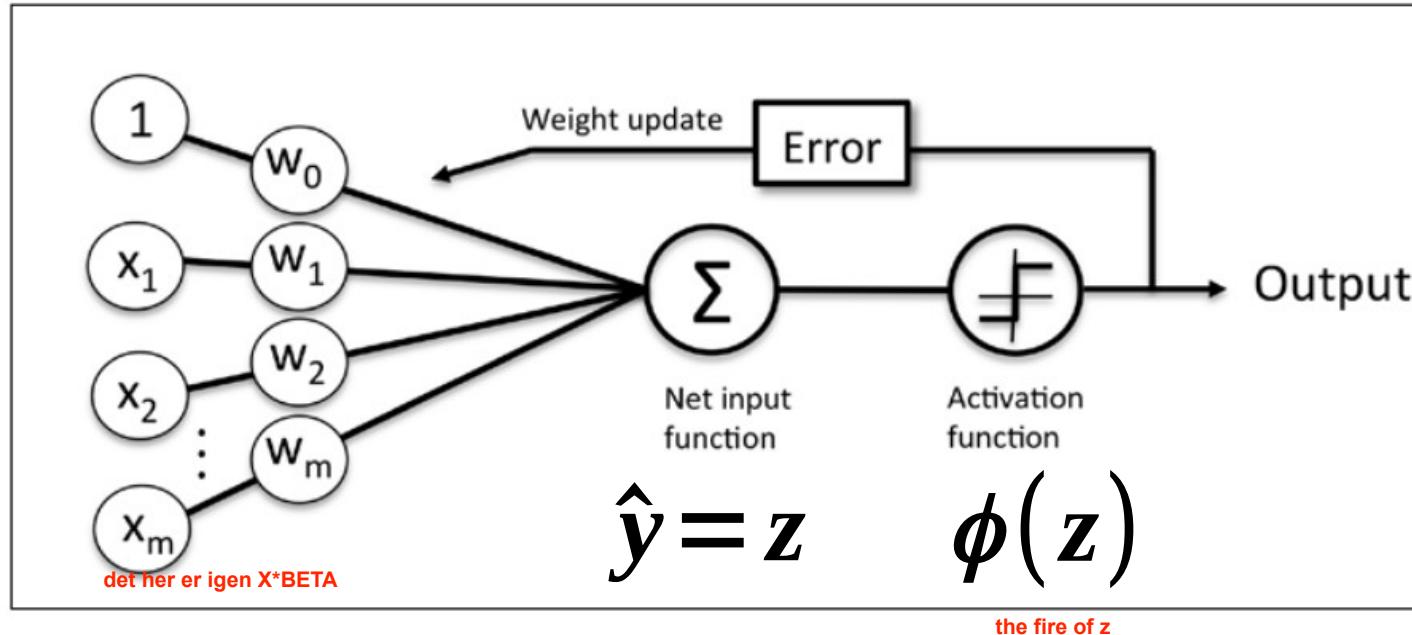
hvis dine variabler er standardiseret, så er 0.1 et godt start punkt for learning rate

Convergence only possible when linearly separable



(p. 23: Raschka, 2015)

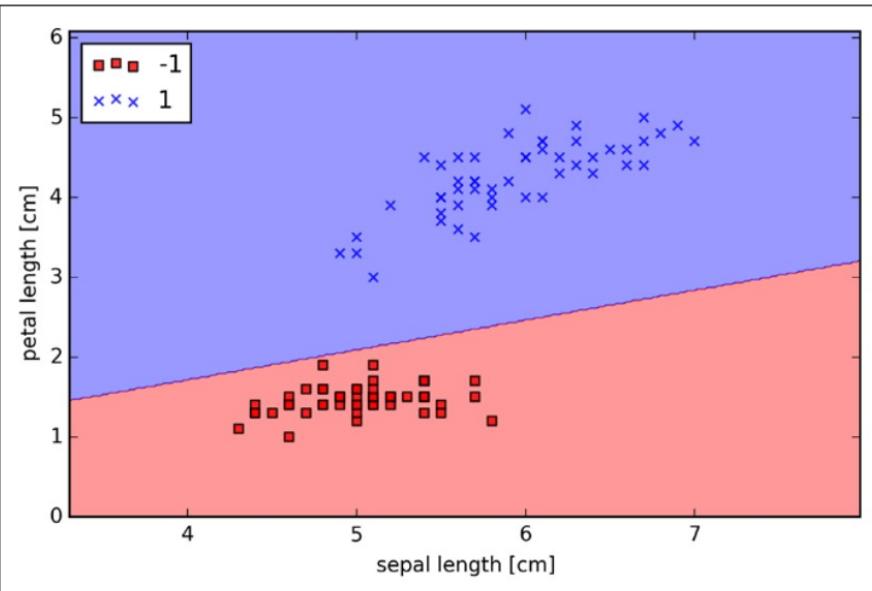
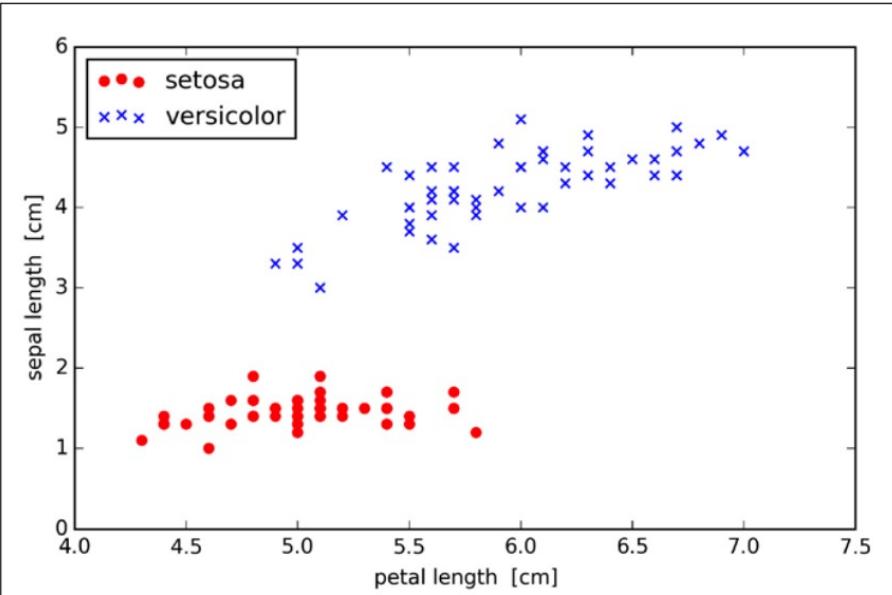
Perceptron: Graphical summary



(p. 24: Raschka, 2015)

An example

the perceptron classifies everything in either blue or red area, depending on sepal and petal length



```
In [4]: ppn.w  
Out[4]: array([-0.4 , -0.68,  1.82])
```

```
In [12]: X[1, :]  
Out[12]: array([4.9, 1.4])
```

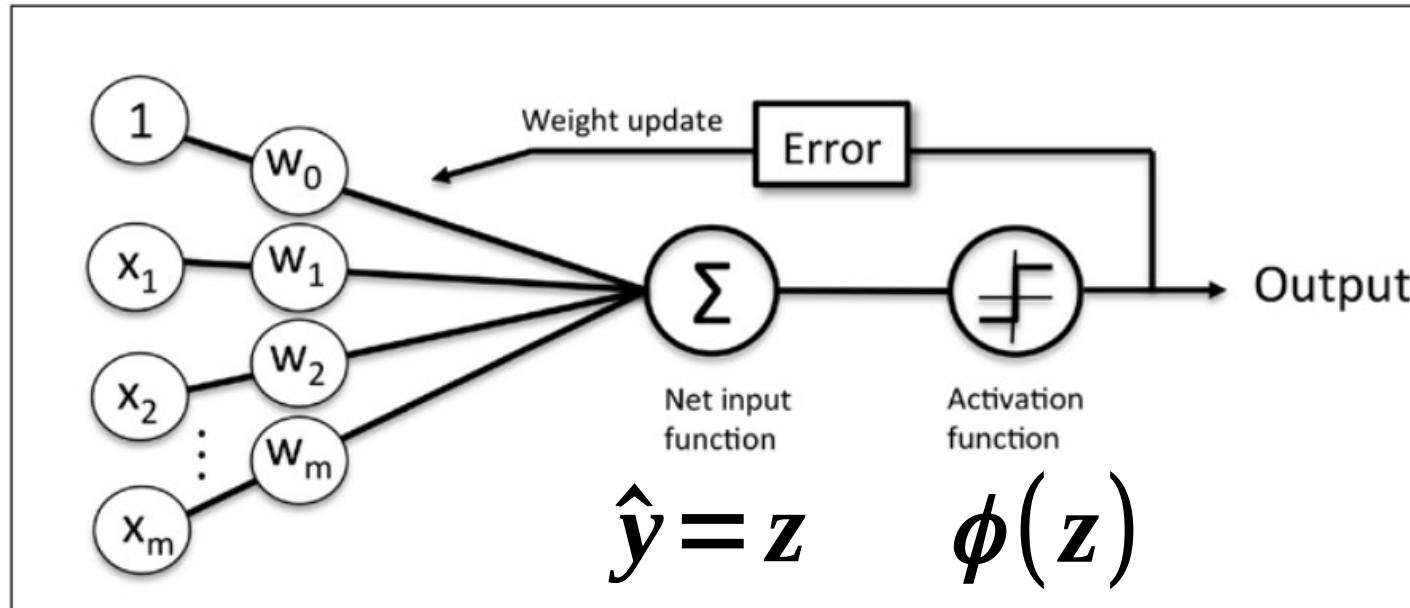
$$\hat{y} = -0.4 - 0.68 \cdot 4.9 + 1.82 \cdot 1.4 = -1.184$$

so its setosa

```
In [11]: ppn.net_input(X[1, :])  
Out[11]: -1.1839999999999975
```

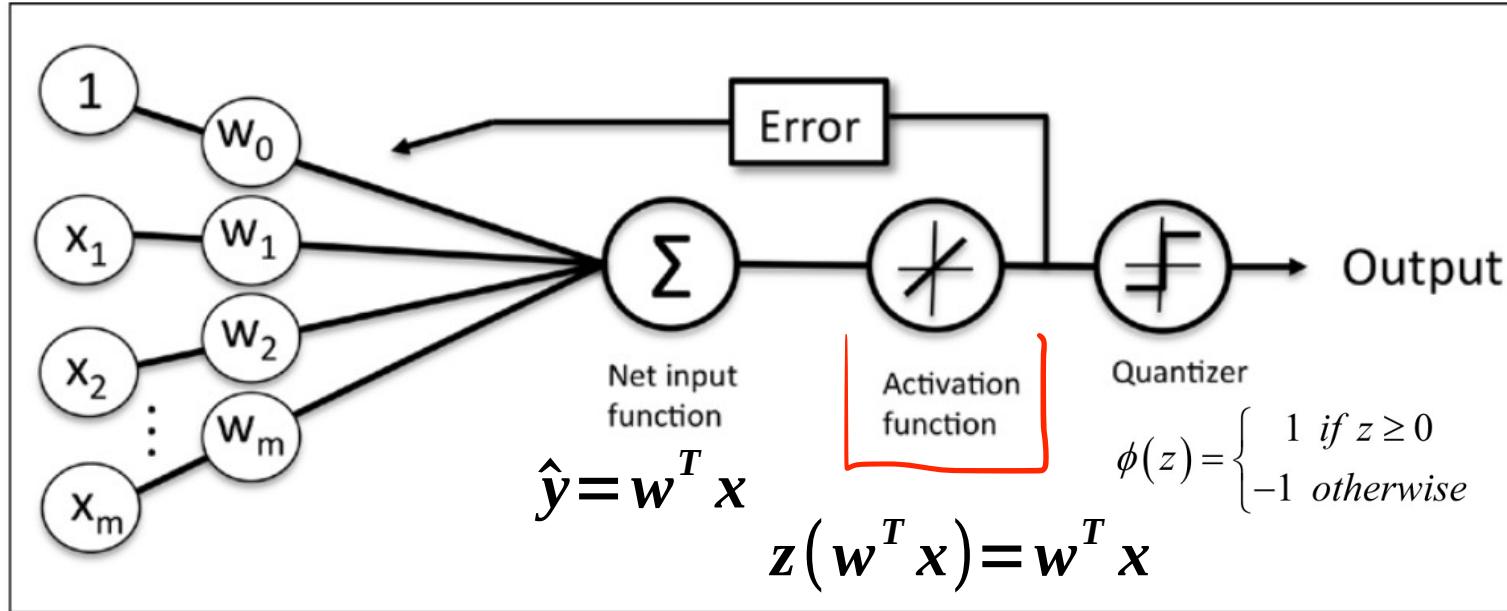
(p. 29 & p. 32: Raschka, 2015)

Perceptron: Graphical summary



(p. 24: Raschka, 2015)

ADAptive LInear NEuron (ADALINE)



$$w^T x = w_0 x_0 + w_1 x_1 + \dots + w_{m-1} x_{m-1} + w_m x_m$$

(p. 33: Raschka, 2015)

ADALINE Gradient descent

```
def __init__(self, eta=0.01, n_iter=50):  
    self.eta = eta  
    self.n_iter = n_iter
```

```
class AdalineGD(object):  
    """ ADaptive LInear NEuron classifier  
  
    Parameters  
    -----  
    eta : float  
        Learning rate (between 0.0 and 1.0)  
    n_iter : int  
        Passes over the training dataset.  
  
    Attributes  
    -----  
    w_ : 1d-array  
        Weights after fitting.  
    errors_ : list  
        Number of misclassifications in every epoch.  
    """
```

Methods

$$\hat{y} = w^T x$$

$$z(w^T x) = w^T x$$

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

```
def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_[1:]) + self.w_[0]
        X*BETA

def activation(self, X):
    """Compute linear activation"""
    return self.net_input(X)

def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.activation(X) >= 0.0, 1, -1)
```

The *fit* method

THE cost function is something we need to minimize!

small constant, so you go slowly down the slope

eta: learning rate (a constant)

output: $X\beta = \hat{y}$

errors: $y - \hat{y}$

$X.T.dot(errors)$: $X^T \cdot (y - \hat{y})$

$X^T \cdot (y - \hat{y}) = \Delta w_1 + \Delta w_2 + \dots + \Delta w_{m-1} + \Delta w_m$

cost function: $(\sum (y - \hat{y})^2) / 2$

```
def fit(self, X, y):
    """ Fit training data.

    Parameters
    -----
    X : {array-like}, shape = [n_samples, n_features]
        Traing vectors, where n_samples
        is the number of samples and
        n_features is the number of features.
    y : array-like, shape = [n_samples]
        Target values.

    Returns
    -----
    self : object

    """
    self.w_ = np.zeros(1 + X.shape[1])
    self.cost_ = []

    for i in range(self.n_iter):
        output = self.net_input(X)
        errors = (y - output)          this gives you updated weights
        self.w_[1:] += self.eta * X.T.dot(errors)
        self.w_[0] += self.eta * errors.sum()
        cost = (errors**2).sum() / 2.0
        self.cost_.append(cost)

    return self
```

cost is sum of errors ^2

then we calculate new errors, are they better or worse?

better? okay lets continue update weights in this direction (down the valley, gradient descent), where the cost function has its global minimum

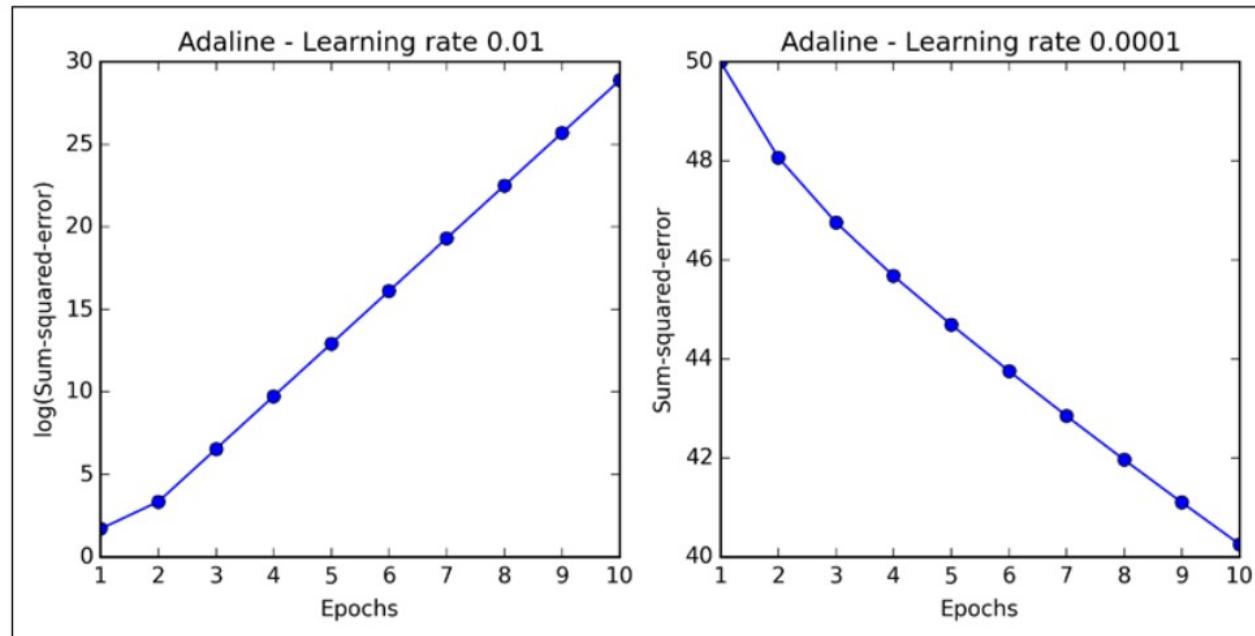
we want to find a weight that minimizes the cost function

Learning rate

Overshooting
the global
minimum

in this case:
high learning rate = error increase (SSR increases)

in this case:
low learning rate = error decreases slowly



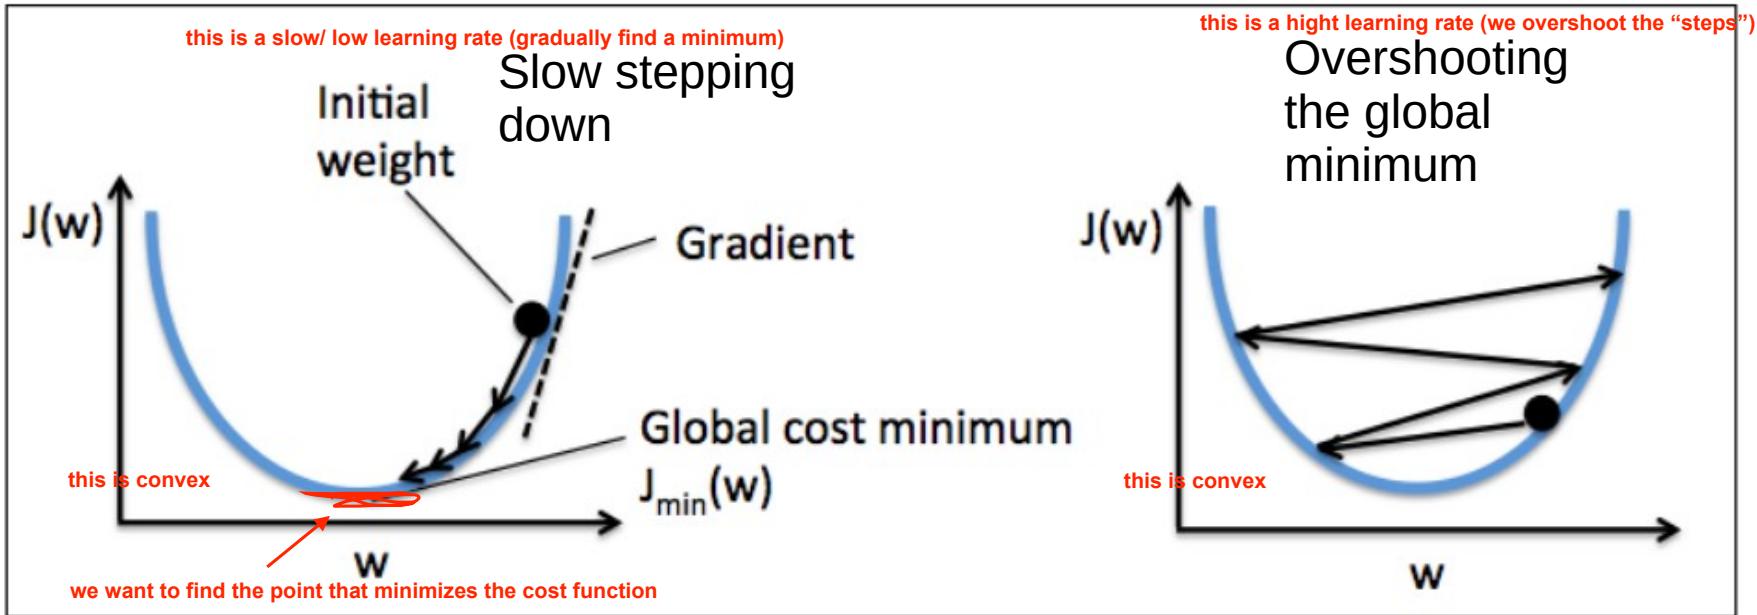
Slow stepping
down

converging would be a straight line, which we could not improve

Always check whether it converges

(p. 40: Raschka, 2015)

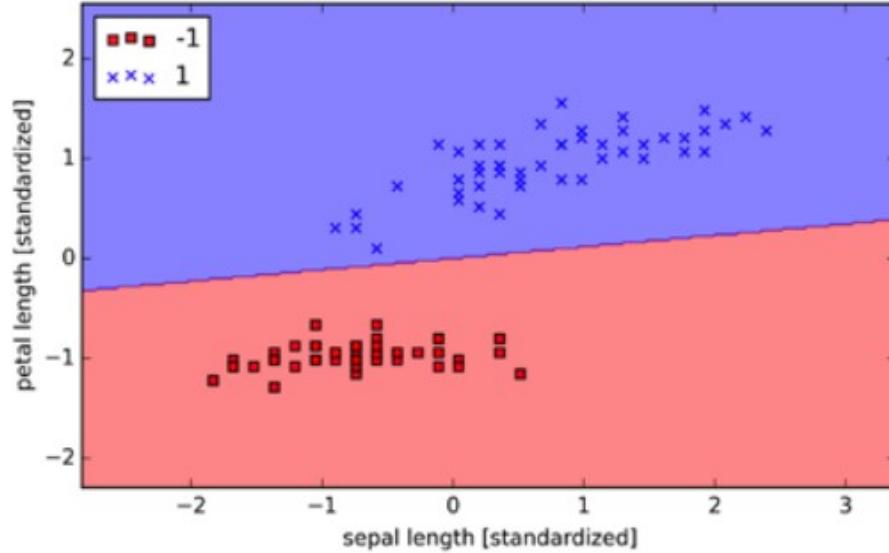
Gradient descent



$$J(w) = \left(\sum (y - \hat{y})^2 \right) / 2$$

(p. 40: Raschka, 2015)

Adaline - Gradient Descent



(p. 42: Raschka, 2015)

```
In [50]: X_std[1, :]  
Out[50]: array([-0.89430898, -1.01435952])
```

```
In [38]: ada.w  
Out[38]: array([ 1.59872116e-16, -1.26256159e-01,  1.10479201e+00])
```

$$\hat{y} \approx 0 - 0.127 \cdot (-0.894) + 1.10 \cdot (-1.014) = -1.01$$



```
In [49]: ada.net_input(X_std[1, :])  
Out[49]: -1.0077442758158444
```

linear activation:
this number -1.01 means its a red dot in the plot

How does this relate to linear regression?

From ADALINE

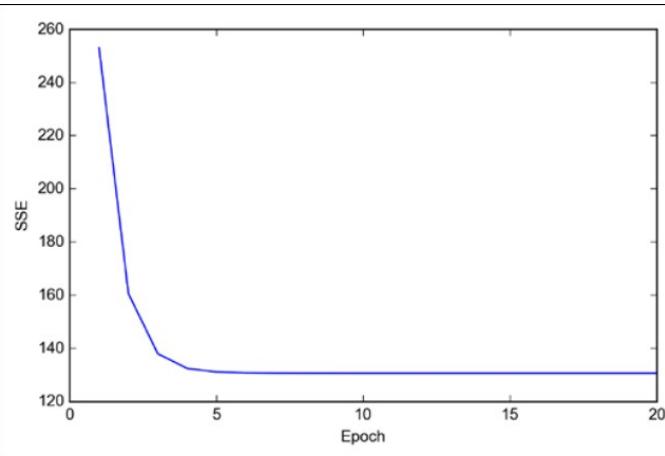
$$\mathbf{w}^T \mathbf{x} = w_0 x_0 + w_1 x_1 + \dots + w_{m-1} x_{m-1} + w_m x_m$$

this time we predict on a continuois scale

Very similar to ADALINE
and when converged will be virtually
identical to the ordinary least
squares solution

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

Convergence



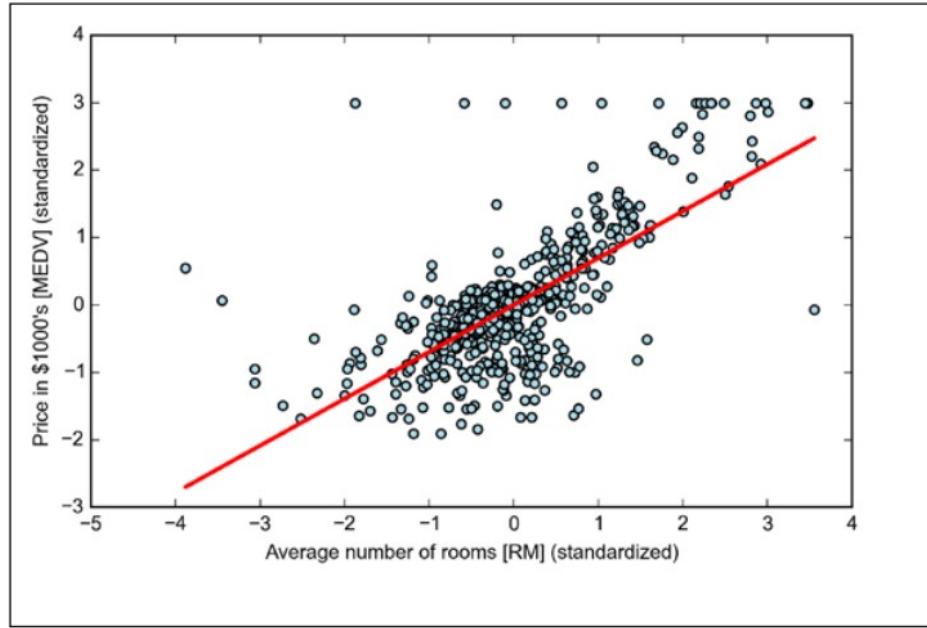
```
class LinearRegressionGD(object):  
  
    def __init__(self, eta=0.001, n_iter=20):  
        self.eta = eta  
        self.n_iter = n_iter  
  
    def fit(self, X, y):  
        self.w_ = np.zeros(1 + X.shape[1])  
        self.cost_ = []  
  
        for i in range(self.n_iter):  
            output = self.net_input(X)  
            errors = (y - output)  
            self.w_[1:] += self.eta * X.T.dot(errors)  
            self.w_[0] += self.eta * errors.sum()  
            cost = (errors**2).sum() / 2.0  
            self.cost_.append(cost)  
        return self  
  
    def net_input(self, X):  
        return np.dot(X, self.w_[1:]) + self.w_[0]  
  
    def predict(self, X):  
        return self.net_input(X)
```

same as adaline, just predicting continous data

What we have seen before

```
lr = LinearRegressionGD()  
lr.fit(X_std, y_std)
```

(p. 288: Raschka, 2015)



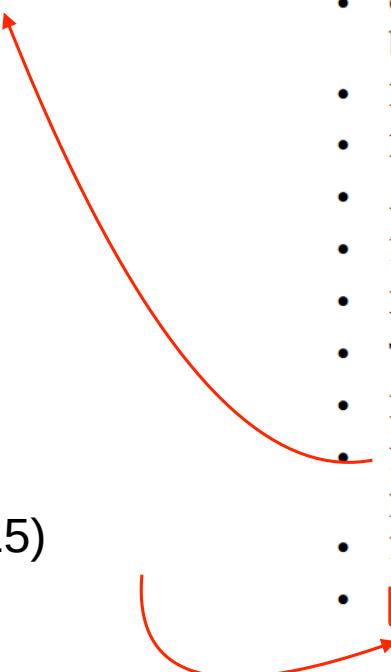
```
In [8]: lr.w_  
Out[8]: array([-4.68958206e-16, 6.95359426e-01])
```

Intercept and slope ↗

Olivier Grisel
@ogrisel

In scikit-learn 1.0, we decided to deprecate the `sklearn.datasets.load_boston` function because the design of this dataset casually assumes that people prefer to buy housing in racially segregated neighborhoods.

(p. 229: Raschka, 2015)



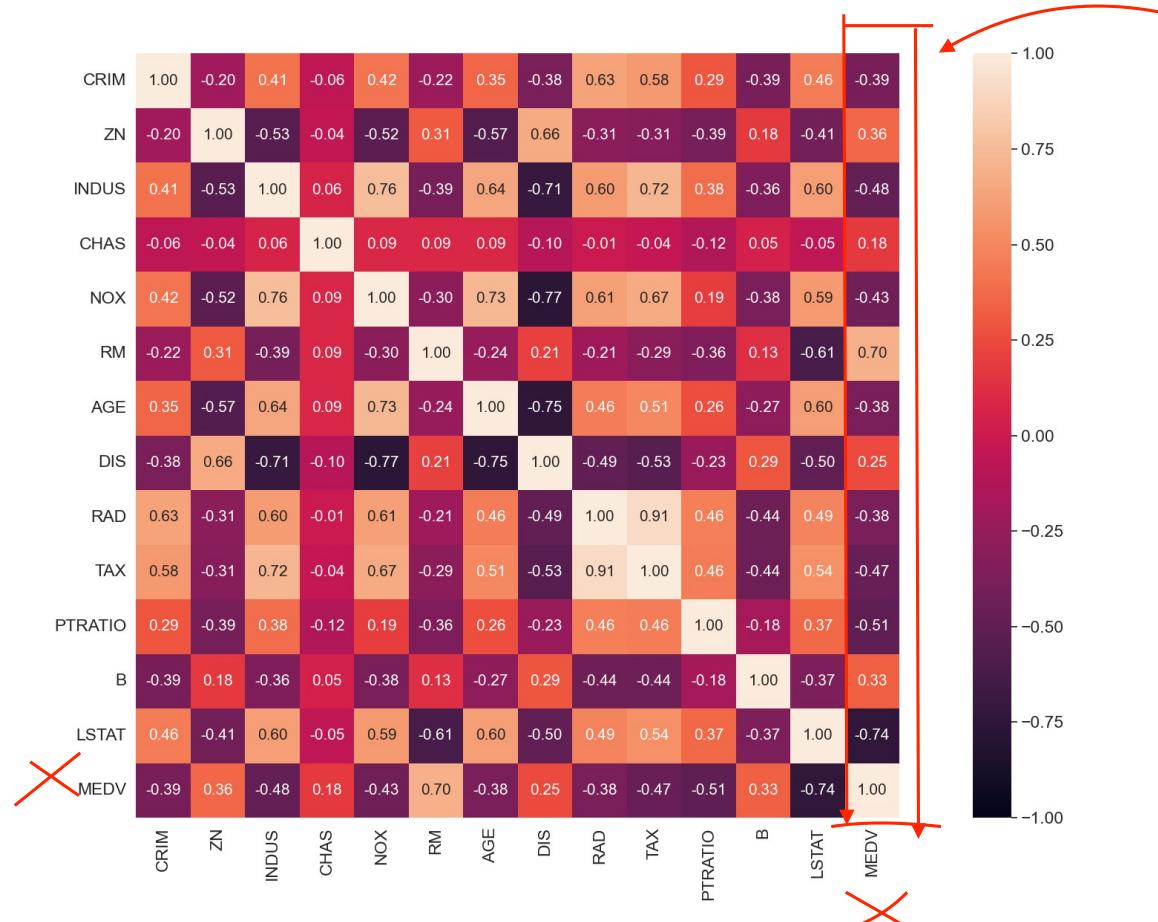
The features of the 506 samples may be summarized as shown in the excerpt of the dataset description:

- **CRIM:** This is the per capita crime rate by town
- **ZN:** This is the proportion of residential land zoned for lots larger than 25,000 sq.ft.
- **INDUS:** This is the proportion of non-retail business acres per town
- **CHAS:** This is the Charles River dummy variable (this is equal to 1 if tract bounds river; 0 otherwise)
- **NOX:** This is the nitric oxides concentration (parts per 10 million)
- **RM:** This is the average number of rooms per dwelling
- **AGE:** This is the proportion of owner-occupied units built prior to 1940
- **DIS:** This is the weighted distances to five Boston employment centers
- **RAD:** This is the index of accessibility to radial highways
- **TAX:** This is the full-value property-tax rate per \$10,000
- **PTRATIO:** This is the pupil-teacher ratio by town
- **B:** This is calculated as $1000(Bk - 0.63)^2$, where Bk is the proportion of people of African American descent by town
- **LSTAT:** This is the percentage lower status of the population
- **MEDV:** This is the median value of owner-occupied homes in \$1000s

Multiple linear regression

$$\hat{MEDV} = \underset{\text{intercept}}{w_0 x_0} + w_1 CRIM + w_2 ZN + w_3 INDUS + w_4 CHAS \\ + w_5 NOX + w_6 RM + w_7 AGE + w_8 DIS \\ + w_9 RAD + w_{10} TAX + w_{11} PTRATIO + w_{12} B + w_{13} LSTAT + \epsilon$$

Collinearity – correlation matrix

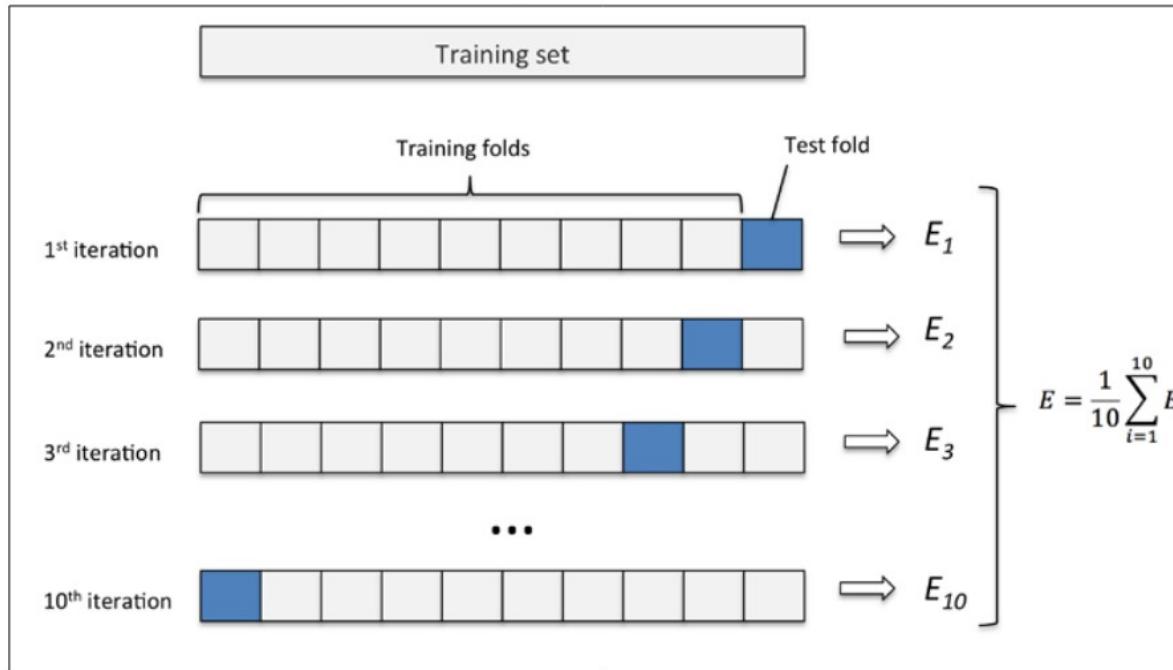


Because of the collinearity, we know we are prone to overfitting, so we do **out-of-sample** prediction instead of validating our model with traditional measures like R^2 and maximum likelihood

we are going to try our model on NEW DATA from today!

How to choose the **out-of-sample** dataset?

Cross-validation

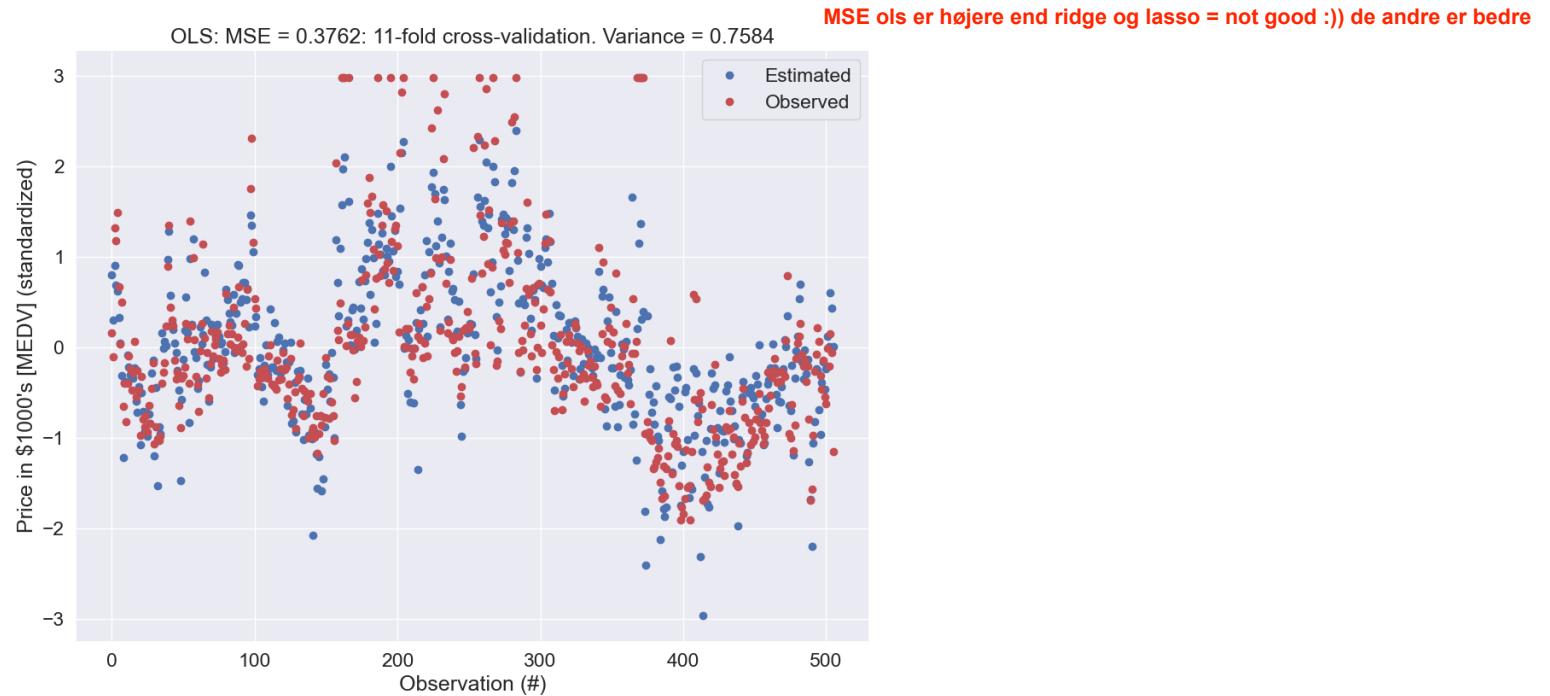


(p. 176: Raschka, 2015)

how to LM in python

```
OLS = LinearRegression()  
OLS.fit(X_std, y_std)  
  
MSE = np.mean(cross_validate(OLS, X_std, y_std, k=11))
```

LAU GØR DET SAMME PÅ SLIDE 38, MEN MERE SIMPELT



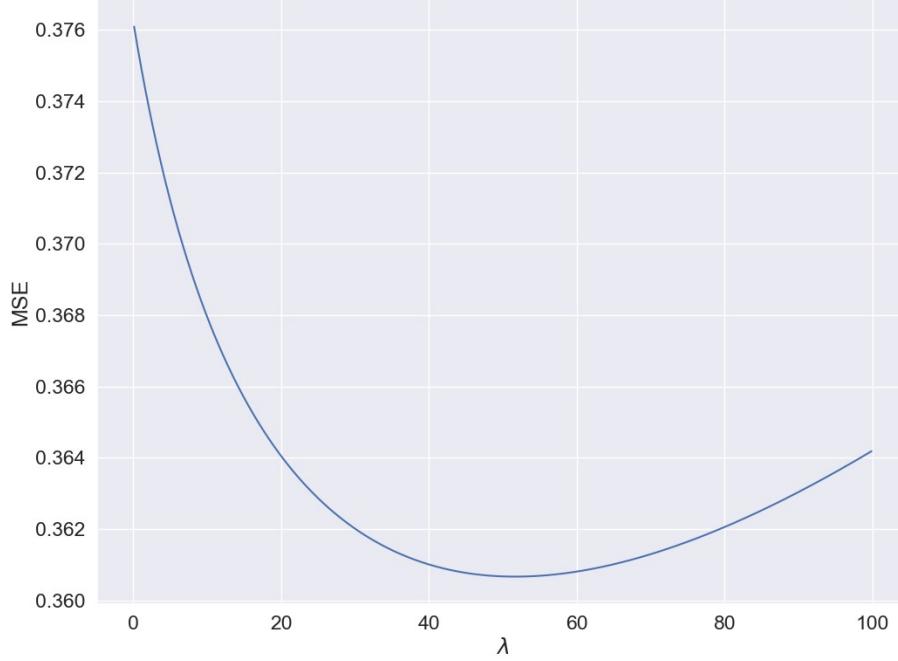
We can impose penalties

(but not on the intercept)

$$J(w)_{Ridge} = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \lambda \|w\|_2^2$$

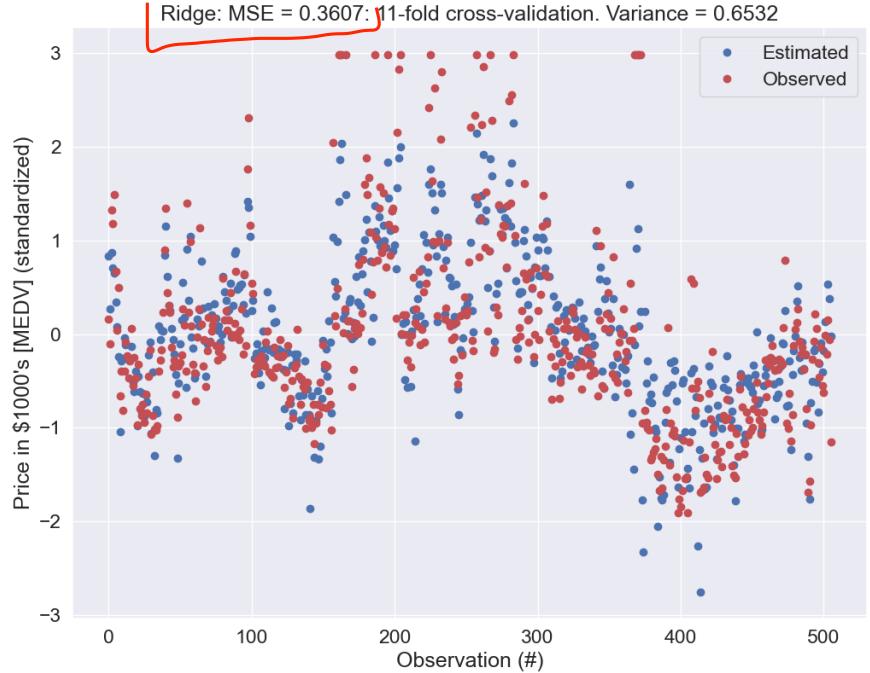
afprøver forskellige lambda values

Ridge Regression: min MSE (0.3607) at: $\lambda = 51.7$



VED AT PUTTE LAMBDA IND!!!!

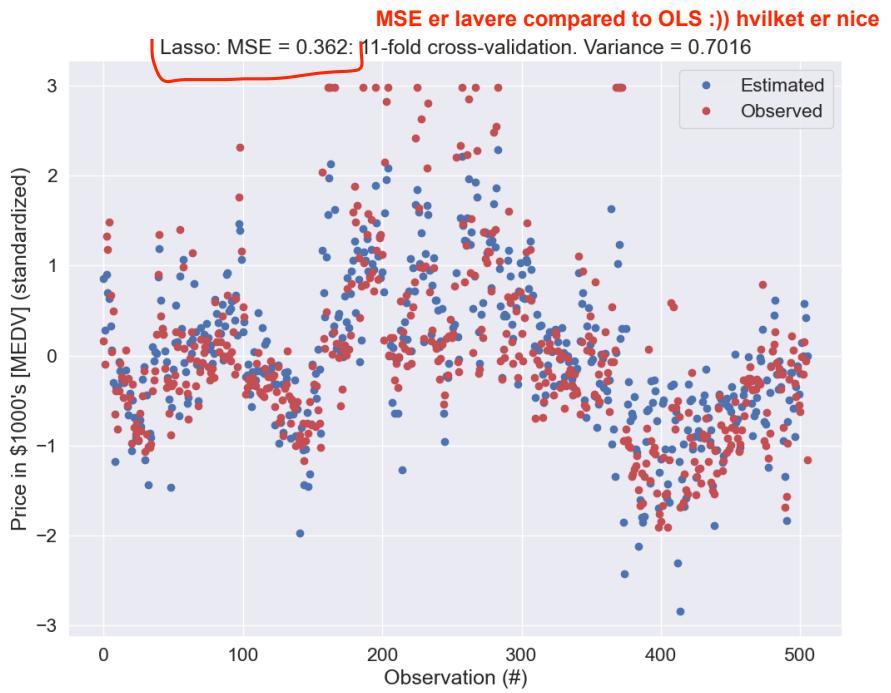
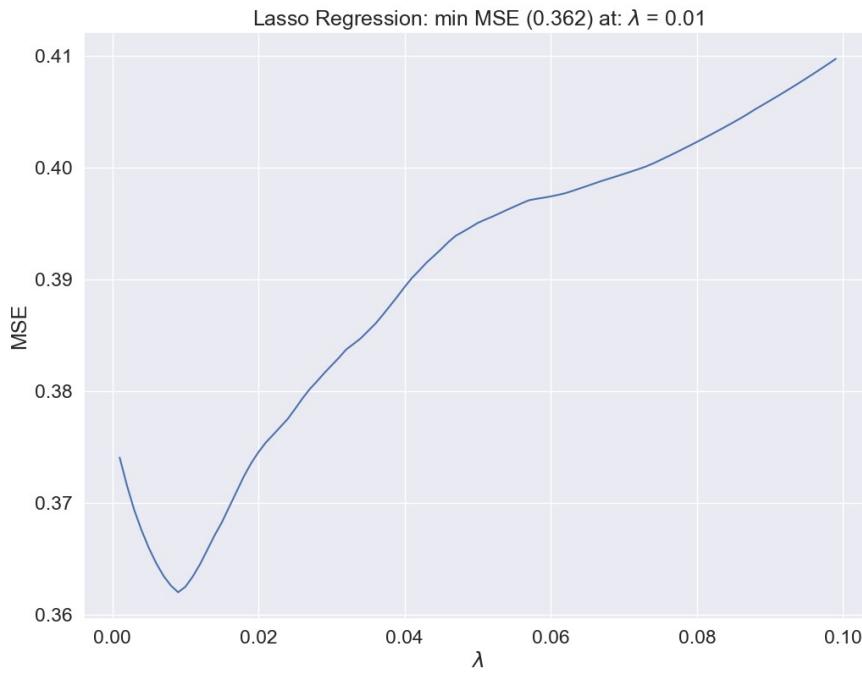
MSE er lavere compared to OLS :)) hvilket er nice



We can impose penalties

(but not on the intercept)

$$J(w)_{LASSO} = \sum_{i=1}^n \left(y^{(i)} - \hat{y}^{(i)} \right)^2 + \lambda \| w \|_1$$



Coefficients are shrunk

```
In [131]: OLS.coef_
Out[131]:
array([-0.09874812,  0.12473758,  0.02386168,  0.06945318, -0.22231612,
       0.2911837 ,  0.0325356 , -0.32907266,  0.34212986, -0.28361575,
      -0.21482076,  0.09763631, -0.43131412])

In [132]: RR.coef_
Out[132]:
array([-0.07536542,  0.08180161, -0.03182673,  0.07855868, -0.13185121,
       0.31082552,  0.00548938, -0.23491485,  0.11242408, -0.0959682 ,
      -0.18436614,  0.09154233, -0.36579723])

In [133]: lasso.coef_
Out[133]:
array([-0.07348948,  0.09107936, -0.          ,  0.07003181, -0.17110318,
       0.30950805,  0.          , -0.29010247,  0.16456993, -0.1319311 ,
      -0.19668957,  0.09063304, -0.41664587])
```

We can impose penalties

(but not on the intercept)

$$J(w)_{ElasticNet} = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \lambda_1 \sum_{j=1}^m w_j^2 + \lambda_2 \sum_{j=1}^m |w_j|$$

don't worry about this, it is more complicated

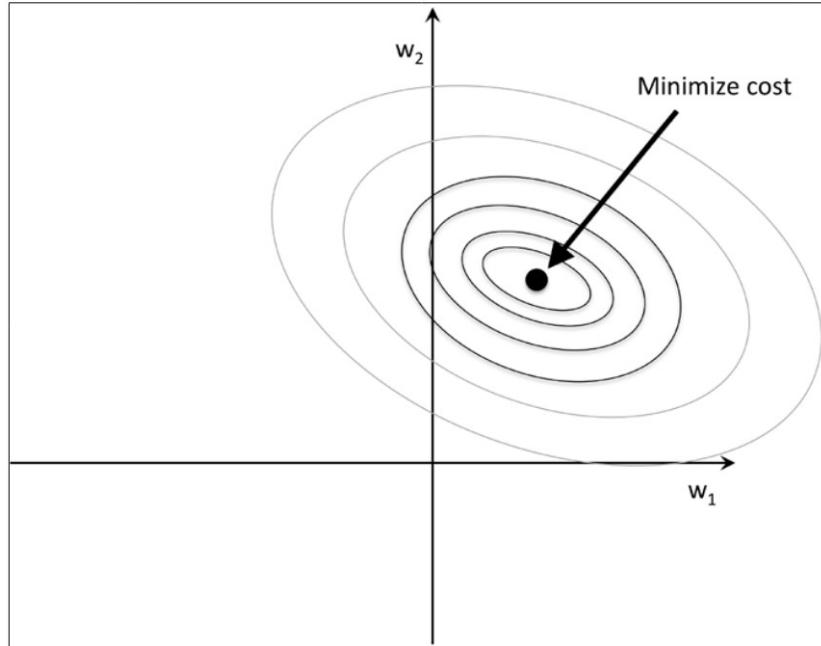
Did you learn?

Linear regression revisited (machine learning)

- 1) Learning some early classification methods
- 2) Learning how linear regression (with biasing penalties) can be constructed and cross-validated
- 3) Understanding that biasing in-sample solutions can improve out-of-sample predictions

Extra slides on regularization

$$J(w) = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



(p. 113: Raschka, 2015)

L2 regularization

Why is the *budget* round?

Compare with a circle centred at (0,0)

$$x^2 + y^2 = r^2$$

$$w_1^2 + w_2^2 = r^2$$

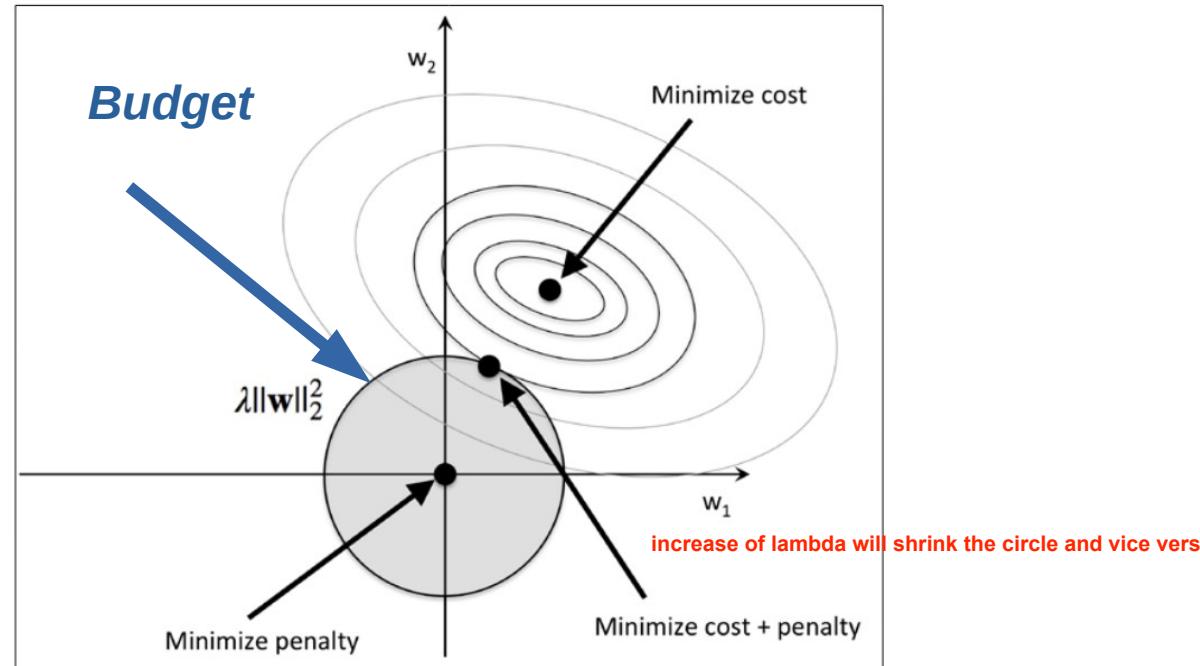
$$\|w\|_2 = \sqrt{w_1^2 + w_2^2}$$

flexibel travel

ecledian norm: the most rational

minimizes the distance the most

(p. 114: Raschka, 2015)



$$J(w)_{Ridge} = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \lambda \|w\|_2^2$$

L1 regularization

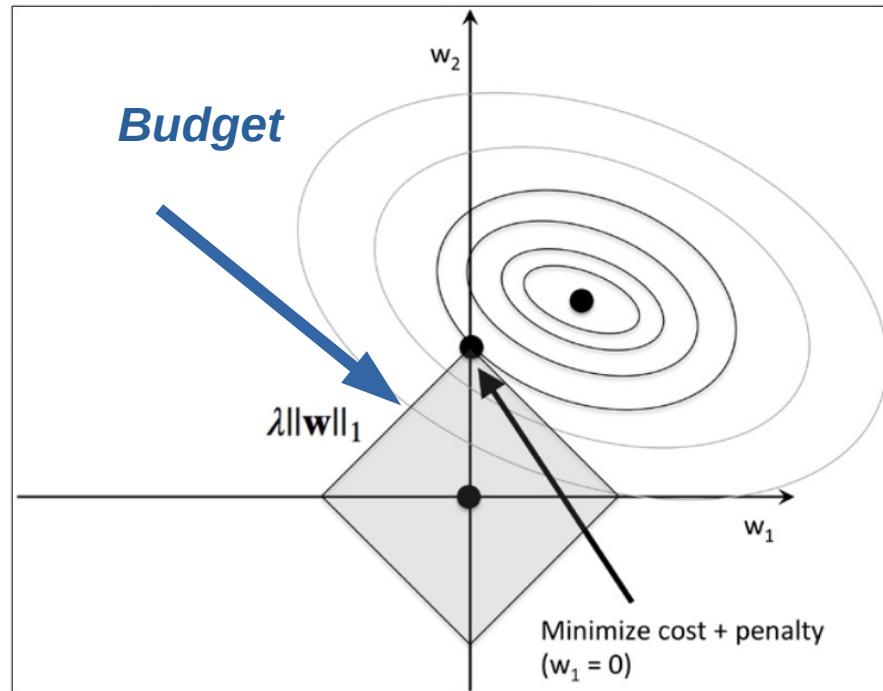
Why is the *budget* square?

manhattan norm(travel), 2 x frem 1 y op fx

$$\|w\|_1 = |w_1| + |w_2|$$

if $w_1 = \max(w_1)$ then $w_2 = 0$

if $w_2 = \max(w_2)$ then $w_1 = 0$



$$J(w)_{LASSO} = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \lambda \|w\|_1$$

References

- Raschka, S., 2015. Python Machine Learning. Packt Publishing Ltd.